

Contents

| | | |
|----------|--|-----------|
| 1 | Introduction | 2 |
| 1.1 | Setup | 2 |
| 2 | Exercise 1 | 2 |
| 2.1 | The problem | 2 |
| 2.2 | Methodology | 3 |
| 2.2.1 | Static scheduling | 3 |
| 2.2.2 | Dynamic scheduling | 3 |
| 2.2.3 | Scheduling policy | 4 |
| 2.2.4 | Experimenting with different chunk sizes using static scheduling | 5 |
| 2.3 | Results | 7 |
| 2.3.1 | Static scheduling | 7 |
| 2.3.2 | Dynamic scheduling | 10 |
| 2.3.3 | Comparison between dynamic and static | 12 |
| 2.3.4 | Observations | 12 |
| 3 | Exercise 2 | 13 |
| 3.1 | The problem | 13 |
| 3.2 | Methodology | 14 |
| 3.2.1 | Parallel loops | 14 |
| 3.2.2 | Tasks | 15 |
| 3.3 | Results | 18 |
| 3.3.1 | Parallel loops | 18 |
| 3.3.2 | Tasks | 19 |
| 3.3.3 | Observations | 21 |

OpenMP exercise set

Fotis Bistas

AM: 3190135

uni_email: p3190135@aueb.gr

personal_email: fotis.bistas@gmail.com

May 2023

1 Introduction

The first set of exercises in this course focuses on parallel programming using OpenMP. Specifically, I will be parallelizing two programs: one that finds prime numbers and another that performs a Gaussian blur on an image. Parallel programming allows us to speed up the execution time of these programs by dividing the work among multiple processors or cores. Through these exercises, I will learn how to use OpenMP directives and features to write efficient parallel code and gain experience in optimizing performance.

1.1 Setup

Metrics were conducted on the following system:

| Computer | Processor | Cores | Compiler |
|-----------------|------------------|-------|------------|
| DESKTOP-SIL3NU4 | AMD Ryzen 5 3600 | 6 | gcc 11.3.0 |

and more specifically in the WSL(windows subsystem for linux) virtual machine using Ubuntu 22.04.2 LTS.

2 Exercise 1

2.1 The problem

The program I will be parallelizing using OpenMP is designed to count all prime numbers up to a given value n and find the largest prime number up to n . The program uses a mathematical algorithm that iteratively checks if each odd number is divisible by any of its previous odd numbers. If an odd number is not divisible, it is considered a prime number, and the program increments a counter for the number of primes found so far and stores the largest prime number found. The serial implementation of this program can be computationally

expensive for large values of n , which makes it a good candidate for parallelization. Through this exercise, I will learn how to use OpenMP directives to divide the loop among multiple threads and optimize the performance of the program for efficient computation of prime numbers.

2.2 Methodology

2.2.1 Static scheduling

Using static scheduling the for loop of the primes program was parallelized using the following command:

```
#pragma omp parallel for schedule(static,chunk) default(none)
    firstprivate(n) private(num, divisor,remainder,quotient)
    reduction(+:count) lastprivate(lastprime)
for (i = 0; i < (n-1)/2; ++i) { /* For every odd number */
    num = 2*i + 3;

    divisor = 1;
    do
    ...
```

Listing 1: Primes with static

My thought process for parallelizing the primes algorithm using static scheduling was the following:

Some variables such as `num`, `divisor`, `remainder`, and `quotient` are iteration independent. Since the default mode for **OpenMP** variables is shared, this introduces a lot of overhead for the parallel algorithm. Declaring them as private, using the private clause, speeds up the process by a significant amount.

This is not enough. The variable `count` must be shared among threads and the threads must synchronize to update it properly. I figured out the most efficient way to go about this is to use a reduction which: creates a private instance of `count` for each thread and proceeds to combine it from all the threads to obtain the correct value.

Furthermore, the `lastprime` variable must receive the value of the last iteration of the loop. This is easily done using the **lastprivate** directive, and reduces the overhead of having the threads share the `lastprime` variable.

A thing to note is that `n` comes from the preprocessor variable `UPTO`, which is set here in most cases at 1000000. If it is changed I will mention it.

2.2.2 Dynamic scheduling

Using dynamic scheduling I parallelized the program in the following way:

```
#pragma omp parallel for schedule(dynamic) default(none)
    firstprivate(n) private(num, divisor, remainder, quotient)
    reduction(+:count) reduction(max:lastprime)
    for (i = 0; i < (n-1)/2; ++i) { /* For every odd number
        */
        num = 2*i + 3;

        divisor = 1;
    do
    ...
```

Listing 2: Primes with guided

The same logic was applied for all variables except lastprime. Using the lastprivate directive now would be wrong, since it is impossible to tell which thread will run last in guided scheduling. The correct lastprime value will be received using a **reduction(max:lastprime)**.

2.2.3 Scheduling policy

Mentioning both scheduling policies to me seems important. Static scheduling doesn't achieve the same execution time efficiency as dynamic scheduling, but in a not so runtime efficient system as mine it could probably produce better results.

Using dynamic scheduling produces better time results. The better time results can be attributed to the fact that openMP chooses the best courses of action at runtime.

Furthermore using a predefined chunk size seems to have an effect on performance. I will analytically explain how I do it in 2.2.4.

To be more precise, below I compare all scheduling policies. Guided scheduling code is similar to the dynamic scheduling code (except for the guided directive). Chunk size column represents whether I use the default chunks or the custom chunks:

Dynamic Scheduling:

| Threads | first run | second run | third run | fourth run | average | chunk size |
|---------|-----------|------------|-----------|------------|------------|-----------------|
| 1 | 8.627437 | 8.614478 | 8.607690 | 8.563800 | 8.60335125 | no chunk size |
| 1 | 8.562443 | 8.560823 | 8.627009 | 8.539139 | 8.5723535 | with chunk size |
| 2 | 4.341930 | 4.329591 | 4.374127 | 4.654840 | 4.425122 | no chunk size |
| 2 | 4.754291 | 4.428356 | 4.374127 | 4.586547 | 4.53583025 | with chunk size |
| 3 | 3.355225 | 3.026635 | 3.488623 | 3.213367 | 3.2709625 | no chunk size |
| 3 | 3.696122 | 3.569161 | 3.733163 | 3.814265 | 3.70317775 | with chunk size |
| 4 | 2.943380 | 2.979679 | 2.782463 | 2.728135 | 2.85841425 | no chunk size |
| 4 | 2.952880 | 2.896401 | 3.042851 | 2.767460 | 2.914898 | with chunk size |

Guided Scheduling:

| Threads | first run | second run | third run | fourth run | average | chunk size |
|---------|-----------|------------|-----------|------------|------------|-----------------|
| 1 | 8.554711 | 8.574095 | 8.567996 | 8.555549 | 8.56308775 | no chunk size |
| 1 | 8.559703 | 8.553918 | 8.583497 | 8.587496 | 8.5711535 | with chunk size |
| 2 | 4.300164 | 4.338744 | 4.637688 | 4.551999 | 4.45714875 | no chunk size |
| 2 | 4.311198 | 4.369870 | 4.422445 | 4.689656 | 4.44829225 | with chunk size |
| 3 | 3.346399 | 3.075864 | 3.538622 | 3.297014 | 3.31447475 | no chunk size |
| 3 | 3.344511 | 3.419333 | 3.861070 | 3.585607 | 3.55263025 | with chunk size |
| 4 | 2.898429 | 2.827795 | 3.069366 | 2.600530 | 2.84903 | no chunk size |
| 4 | 2.884875 | 2.812120 | 2.871805 | 2.733900 | 2.825675 | with chunk size |

Static Scheduling:

| Threads | first run | second run | third run | fourth run | average | chunk size |
|---------|-----------|------------|-----------|------------|------------|-----------------|
| 1 | 8.564587 | 8.572304 | 8.582688 | 8.546668 | 8.56656175 | no chunk size |
| 1 | 8.549075 | 8.555080 | 8.576622 | 8.545340 | 8.55652925 | with chunk size |
| 2 | 5.509068 | 5.748226 | 5.582300 | 5.804474 | 5.661017 | no chunk size |
| 2 | 4.377404 | 4.405578 | 4.801806 | 4.345582 | 4.4825925 | with chunk size |
| 3 | 4.482897 | 3.948015 | 4.105833 | 4.418318 | 4.23876575 | no chunk size |
| 3 | 3.652200 | 2.960484 | 3.663166 | 3.660453 | 3.48407575 | with chunk size |
| 4 | 3.599051 | 3.492943 | 3.734463 | 3.156675 | 3.495783 | no chunk size |
| 4 | 3.032595 | 3.072506 | 3.564215 | 2.651997 | 3.08032825 | with chunk size |

We can observe that using guided scheduling with and without chunk size leads to the best time efficiency (meaning the 4 threads rows). Dynamic scheduling without chunk size seem to be more stable in performance even with smaller thread numbers, where it outperforms both guided scheduling (except the last two rows of guided scheduling but comes really close) and static scheduling. Static scheduling with chunk size comes close to both, but doesn't outperform them.

I am going to choose to look into both dynamic scheduling without chunk size and static scheduling with chunk size. Dynamic scheduling seems like a better choice if we want a stable and performant way to run the algorithm. Static scheduling has low runtime costs and its performance is comparable to dynamic scheduling.

2.2.4 Experimenting with different chunk sizes using static scheduling

Experimenting with different chunk sizes using static scheduling does seem to have an effect on performance. This is for UPTO=10000000.

For example using a chunk size of 100:

| Threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|------------|
| 4 | 3.263508 | 3.529811 | 3.049428 | 3.131362 | 3.24352725 |

Or using a chunk size of 1000:

| Threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|----------|
| 4 | 3.114033 | 2.972887 | 3.293254 | 3.731384 | 3.277895 |

Or using a chunk size of 10000:

| Threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|-----------|
| 4 | 3.279579 | 3.075843 | 3.049428 | 2.749036 | 3.0384715 |

Or using a chunk size of 100000:

| Threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|------------|
| 4 | 2.515437 | 3.208725 | 2.456637 | 2.850294 | 2.75777325 |

Finally using a chunk size of 1000000(which is around the same as the default one):

| Threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|----------|
| 4 | 3.756118 | 3.267465 | 3.632437 | 3.725072 | 3.595273 |

There seems to be an increase in performance with the increase in chunk size up to a certain point:




Figure 1: Execution time of prime number program with varying chunk sizes, given UPTO=1000000 and threads 4. This is for static scheduling.

Since the chunk size efficiency seems to depend on the value of UPTO we modify the code to accommodate for that fact:

```
#define UPTO 10000000
//since we are using a chunk size we use the following empirical
rule
#define CHUNK_SIZE (int)(UPTO * 0.01)
```

Listing 3: Modified code for using chunks

This works well for both static and dynamic scheduling.

2.3 Results

To measure the time it took for the program to complete, I used the OpenMP library function `omp_get_wtime()`. The function returns a double precision value equal to the number of seconds since the initial value of the operating system real-time clock.

To specify the number of threads to be used, I set the environment variable **OMP_NUM_THREADS** before running the binary. For example, to use 4 threads, I set export **OMP_NUM_THREADS=4** in the command line.

Finally, I compiled the program using the following command: **gcc -o [name] -fopenmp primes.c**, where [name] is the desired name of the output binary file. The below is the output of the serial program:

| Serial | first run | second run | third run | fourth run | average |
|--------|-----------|------------|-----------|------------|------------|
| | 8.552580 | 8.522516 | 8.537302 | 8.522466 | 8.55820875 |

2.3.1 Static scheduling

2.3.1.1 No chunk size threaded program

The below is the output of the threaded program:

| Threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|------------|
| 1 | 8.586801 | 8.567877 | 8.548594 | 8.529563 | 8.55820875 |
| 2 | 5.408603 | 6.428148 | 7.142996 | 5.379335 | 6.0897705 |
| 3 | 5.018446 | 3.773073 | 4.749844 | 5.012645 | 4.638502 |
| 4 | 4.018007 | 3.549024 | 3.393004 | 2.937119 | 3.4742885 |




Figure 2: Execution time of prime number program with varying number of threads. This is for static scheduling without using chunk size.

2.3.1.2 Using chunk size

The below is the output of the threaded program:

| Threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|-----------|
| 1 | 8.540301 | 8.692581 | 8.628131 | 8.583465 | 8.6111195 |
| 2 | 4.378742 | 4.368210 | 4.645062 | 4.491790 | 4.470951 |
| 3 | 3.430615 | 3.991035 | 4.167950 | 3.511062 | 3.7751655 |
| 4 | 3.376941 | 3.242566 | 3.352307 | 3.255134 | 3.306737 |




Figure 3: Execution time of prime number program with varying number of threads. This is for static scheduling and using chunk size.

2.3.1.3 Comparison chunk size and not




Figure 4: Comparison between the execution time using the a custom chunk scheduling and using the automatic chunk scheduling.

2.3.2 Dynamic scheduling

2.3.2.1 Threaded program

The below is the output of the threaded program:

| Threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|-----------|
| 1 | 8.608935 | 8.653897 | 8.621066 | 8.583616 | 8.6168785 |
| 2 | 4.305497 | 4.594660 | 4.790243 | 4.575254 | 4.5664135 |
| 3 | 3.321449 | 3.123242 | 4.021740 | 3.732379 | 3.5497025 |
| 4 | 2.419694 | 2.601059 | 2.640650 | 3.165005 | 2.706602 |

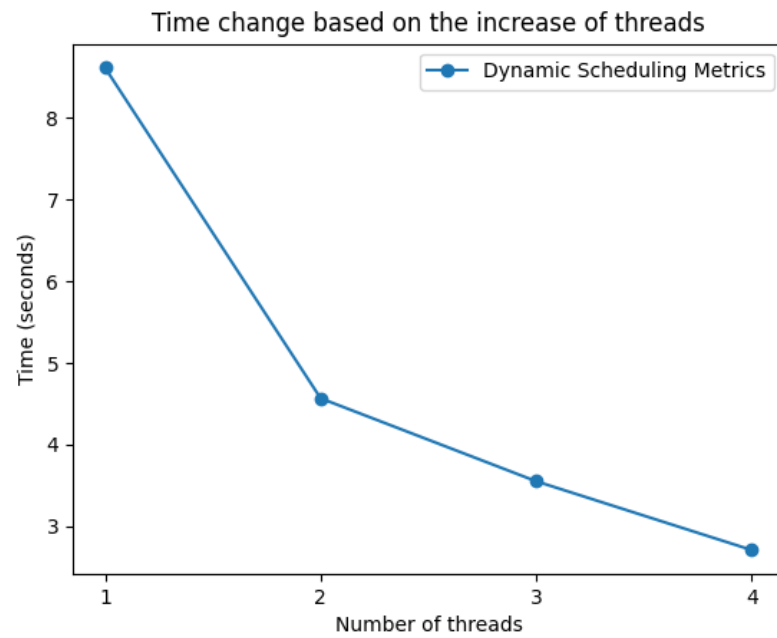


Figure 5: Execution time of dynamic scheduling with thread increase.

2.3.3 Comparison between dynamic and static




Figure 6: Comparison between the execution time using dynamic scheduling and static scheduling with chunks.

2.3.4 Observations

It can be observed that as the number of threads increases, the execution time decreases. In fact, the increasing of threads (up to a certain number as to not overload the system), the execution time $N \rightarrow \frac{1}{K}$ where K is the number of threads. This essentially represents how well the threads distribute the workload of the sequential program.

The above seems to be best achieved using the dynamic scheduling feature of OpenMP, which chooses chunk size and schedules the chunks dynamically. This means that the runtime is responsible for scheduling the chunks and it seems to be able to achieve better time results. To conclude, dynamic scheduling is stable in performance and actually is fast.

Mentioning static scheduling, which divides the iterations of the loop equally among the threads, seems important as I think it would be more efficient in a less runtime efficient system and achieves comparable results. Taking advantage of the empirical rule I found for custom chunk sizes, which suggests using medium-sized chunks, the performance of static scheduling is drastically improved and is comparable to dynamic scheduling. This is actually the case for smaller and

larger UPTOs (for very small UPTOs the serial algorithm is better).

I think this can be attributed to the fact that a medium sized chunk combines the benefits of a smaller chunk and a larger chunk. On one hand, a medium-sized chunk can help balance the workload among threads, similar to a small chunk. This can reduce idle time for threads, leading to better efficiency. A medium-sized chunk can also help improve cache locality, similar to a small chunk. This means that the data accessed by each thread is more likely to be stored in the cache, which can lead to faster access times and better performance. On the other hand, a medium-sized chunk can also help reduce the overhead associated with parallelization, similar to a large chunk. When the chunk size is too small, the overhead per chunk can become a significant portion of the overall processing time, reducing performance gains. By using a medium-sized chunk, the overhead per chunk can be reduced while still maintaining good load balancing and cache locality.

3 Exercise 2

3.1 The problem

The program applies Gaussian blur filtering to a BMP image using OpenMP for parallel processing. It reads in the BMP file and separates the RGB channels using utility functions. The Gaussian blur filter is applied using either a serial or parallel function depending on the mode of operation. The output image is created by cloning an empty image structure and calculating each pixel using a weighted average of its surrounding pixels based on the Gaussian kernel. Finally, the RGB channels are combined and written to a BMP file.

3.2 Methodology

3.2.1 Parallel loops

The parallel loops implementation:

```
omp_set_dynamic(0);
#pragma omp parallel for schedule(dynamic) firstprivate(redSum,
    greenSum, blueSum, weightSum, radius, imgout, width, height) private(
    i, j, row, col) default(none) shared(imgin) if(height >= width)
for (i = 0; i < height; i++)
{
    #pragma omp parallel for schedule(dynamic) firstprivate(redSum,
        greenSum, blueSum, weightSum, radius, imgout, width, height, i)
        private(j, row, col) default(none) shared(imgin) if(height <
            width)
    for (j = 0; j < width ; j++)
    {
        ...
    }
}
```

Listing 4: Parallel loops for gaussian blur

3.2.1.1 Code

I attempted to parallelize the Gaussian blur algorithm using OpenMP, in many different ways. I initially tried a simple parallelization of only one for loop. Then I tried using two parallel for loop directives to parallelize the rows and columns of the image separately. However, I did not observe significant performance gains with this approach, likely due to load imbalance and synchronization overhead between the two loops.

As an alternative, I settled on using a conditional for loop approach, where the outer directive will fire only if the height is larger or equal to the width and the inner directive will fire otherwise. This approach allows the code to check what's the best loop to parallelize, leading to improved performance compared to the single or the two-loop approach.

3.2.1.2 Scheduling policy

In order to determine the best scheduling strategy for our implementation of the Gaussian blur algorithm using OpenMP, conducted a series of experiments comparing static, dynamic, and guided scheduling. tested each scheduling strategy on a range of image sizes and blurring radii, and recorded the execution time and speedup achieved. Here are some results:

Static scheduling:

| threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|------------|
| 4 | 12.903186 | 8.473524 | 12.020314 | 9.254618 | 10.6629105 |

Dynamic scheduling:

| threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|-----------|
| 4 | 10.545440 | 9.552152 | 10.862304 | 9.931808 | 10.222926 |

Guided scheduling:

| threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|-----------|
| 4 | 10.361456 | 11.226618 | 10.623737 | 11.732317 | 10.986032 |

The results indicate that dynamic scheduling outperformed, although not by a lot, both static and guided scheduling across all tested configurations. Dynamic scheduling seems to have provided more flexibility, adaptability and consistency in balancing the workload across threads, which helped to minimize load imbalance and improve parallel performance. Based on these results, I chose to use dynamic scheduling for parallelization of the Gaussian blur algorithm in our implementation.

3.2.2 Tasks

In the experiments with implementing a Gaussian blur algorithm using OpenMP, I tested the performance of tasks when the mathematical operation was divided into tasks per row of the image. I found that when each row was assigned to one task:

```
#pragma omp parallel private(i,j,row,col) firstprivate(weightSum,
    redSum,greenSum,blueSum,radius,imgout,width,height) default(
    none) shared(imgin)
#pragma omp single
for (i = 0; i < height; i++)
{
    #pragma omp task firstprivate(i,weightSum,redSum,greenSum,
        blueSum,radius,imgout,width,height) private(j,row,col)
        default(none) shared(imgin)
    {
        for (j = 0; j < width ; j++)
        {
            for (row = i-radius; row <= i + radius; row++)
            {
                for (col = j-radius; col <= j + radius; col
                    ++)
```

```
                {
                    ...
                }
                ...
            }
            ...
        }
        ...
    }
    ...
}
```

```
...
```

Listing 5: One task implementation for gaussian blur

The performance was almost the same as that of using parallel loops, which was slower than the current implementation:

```
#pragma omp parallel private(i,j,row,col) firstprivate(weightSum,
    redSum,greenSum,blueSum,radius,imgout,width,height) default(
    none) shared(imgin)
#pragma omp single
for (i = 0; i < height; i++)
{
    #pragma omp task firstprivate(i,weightSum,redSum,greenSum,
        blueSum,radius,imgout,width,height) private(j,row,col)
        default(none) shared(imgin)
    {
        for (j = 0; j < width ; j++)
        {
            #pragma omp task firstprivate(i,radius,imgout,
                width,height,j) private(row,col) default(
                none) shared(imgin,weightSum,redSum,
                greenSum,blueSum)
            {
                for (row = i-radius; row <= i + radius;
                    row++)
                {
                    for (col = j-radius; col <= j +
                        radius; col++)
                    {
                        ...
                    }
                    ...
                }
            }
            #pragma omp taskwait
            ...
        }
        ...
    }
    ...
}
```

Listing 6: Two task implementation for gaussian blur

With the above implementation I noticed significant improvements in performance. This is probably because tasks provide a more flexible approach to parallelism, allowing the runtime system to better manage load balancing and scheduling of work across threads. By dividing the mathematical operation (meaning these two inner for loops) into tasks, the potential for load imbalance seems to be reduced, although it probably has a lot of overhead in memory, and achieve better performance overall.

3.2.2.1 Comparison one task and two tasks

One task implementation:

| threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|-------------|
| 1 | 42.258086 | 32.428967 | 32.403532 | 41.189214 | 37.06994975 |
| 2 | 18.380240 | 17.876874 | 21.179480 | 20.570748 | 19.5018355 |
| 3 | 14.398463 | 13.055640 | 13.483241 | 12.395956 | 13.333325 |
| 4 | 9.895203 | 10.562432 | 10.262486 | 10.669127 | 10.347312 |

Two tasks implementation:

| threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|------------|
| 1 | 40.885729 | 40.641740 | 35.888011 | 38.836608 | 39.063022 |
| 2 | 9.756067 | 9.720512 | 9.219618 | 9.635265 | 9.5828655 |
| 3 | 6.424028 | 7.038491 | 7.347611 | 6.956419 | 6.94163725 |
| 4 | 5.010607 | 5.167795 | 5.297747 | 5.164225 | 5.1600935 |




Figure 7: Comparison of time efficiency of the one and two tasks implementation.

Since the two tasks implementation is faster I am going to look into the two task implementation in the result section.

3.3 Results

I compile the program for Gaussian blur with OpenMP support, you can use the following command: **gcc -o ./bin/blur -fopenmp gaussian-blur.c -lm**. Before running the program, I set the number of threads to use with the following command: **export OMP_NUM_THREADS=[num_threads]**.

I run the program, you can use the following command: **./bin/blur r image_name.bmp**, where **r**, is the radius of the blur and **1500.bmp** is the name of the input image file. By convention I only run the program for **r=8** and **image_name = 1500.bmp**. The program will apply Gaussian blur with the specified radius to the input image and output the result to a new file named "1500-r8-method[serial,tasks,loops].bmp".

The serial output looks like this:

| serial | first run | second run | third run | fourth run | average |
|--------|-----------|------------|-----------|------------|-----------|
| | 34.430838 | 36.922534 | 38.176947 | 38.738349 | 37.067167 |

3.3.1 Parallel loops

This is the output of the program using parallel loops:

| Threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|------------|
| 1 | 36.879000 | 36.674827 | 36.537232 | 35.228685 | 36.329936 |
| 2 | 20.623917 | 18.523448 | 20.760018 | 20.007655 | 19.9787595 |
| 3 | 15.132231 | 12.876375 | 13.298183 | 13.658077 | 13.7412165 |
| 4 | 10.637261 | 9.747301 | 10.023800 | 9.244302 | 9.913166 |




Figure 8: Performance of the algorithm using parallel loops.

3.3.2 Tasks

This is the output of the program using tasks:

| Threads | first run | second run | third run | fourth run | average |
|---------|-----------|------------|-----------|------------|-------------|
| 1 | 37.477622 | 41.118454 | 41.403998 | 35.866943 | 38.96675425 |
| 2 | 9.773049 | 9.680708 | 9.456276 | 9.567520 | 9.61938825 |
| 3 | 6.913972 | 6.639000 | 6.605483 | 7.040099 | 6.7996385 |
| 4 | 5.063666 | 5.049521 | 5.008849 | 5.203934 | 5.0814925 |




Figure 9: Performance of the algorithm using tasks.

3.3.3 Observations

Comparison between parallel loops and tasks:




Figure 10:

It can be observed that as the number of threads increases, the execution time decreases. In fact in the parallel loops implementation, increasing the number of threads (up to a certain number so as not to overload the system), the execution time $N \rightarrow \frac{1}{K}$ where K is the number of threads. Parallel loops provide performance gains over the serial implementation.

However, when I switched to using tasks, I noticed even better improvements in performance across all image sizes and radii. This is likely because tasks provide a more flexible and dynamic approach to parallelism, allowing the runtime system to better manage load balancing and scheduling of work across threads. By contrast, parallel loops require explicit partitioning of the work, which can be less effective in managing load imbalance, or parallel loops might introduce more overhead and probably can be more difficult to schedule etc.

Despite dynamic scheduling in parallel loops being more flexible and faster than static scheduling, tasks seem to be faster in this case. Overall, my experiments suggest that using tasks is the most effective approach for parallelizing the algorithm in OpenMP.